

PLATFORM EVALUATION

Lyzr

Agent Studio

Build, Deploy & Manage AI Agents — Feature Overview

What is Lyzr?

Lyzr is an enterprise-grade AI agent platform that lets developers and business users build, deploy, and manage intelligent AI agents.

Developers

Python ADK & API-first tooling

Non-Technical

Drag-and-drop Visual Builder

Enterprise

On-prem, guardrails & compliance

Core Promise

Build AI agents with minimal infrastructure overhead

Deployment

Cloud (Lyzr-hosted) or Private / On-Premises

Access

API, Python ADK, Visual Builder & Blueprints

Focus

Responsible AI — Hallucination & Bias controls built-in

Pricing & Plans

Community

\$0/mo

- 10 Agents
- 5 credits/mo
- 5 Knowledge Bases
- 100 MB RAG Storage
- 1 Builder License
- Limited Models

Starter

\$19/mo

- 15 Agents
- 20 credits/mo
- 10 Knowledge Bases
- 100 MB RAG Storage
- 1 Builder License
- Orchestration

Pro

\$99/mo

- 25 Agents
- 100 credits/mo
- 15 Knowledge Bases
- 1 GB RAG Storage
- Super Agents
- All Blueprints

Enterprise

Custom

- Unlimited Agents
- BYOM + All Models
- On-Prem Deploy
- Team Collaboration
- Agent Eval
- Responsible AI

Agent Building Options

Multiple ways to build

Visual Builder

No-code

Drag-and-drop interface for non-technical users to create, configure and deploy agents without writing any code.

Lyzr ADK

Python Library

Developer-first Python SDK for building, deploying, and managing AI agents programmatically with full control.

Lyzr Architect

Text-to-App

Describe what you want in plain text and Lyzr auto-generates the agent app structure for you.

Workflow Builder

Flow Design

Visually design the agent's decision flow — users define the logic, conditions, and branching paths.

Proxy Agents

External Frameworks

Use Lyzr as a control panel for agents built on external AI frameworks like LangChain, CrewAI etc.

Agent Types & Capabilities

Voice Agents

Build agents that interact via voice for real-time conversational experiences.

GitAgent

Turn any Git repository into an intelligent agent that understands your codebase.

LyzrGPT

Built-in conversational AI layer powering chat-based agent interactions.

Skills

Internal abilities agents can execute — search, summarize, query, transform data.

Enterprise Integration

Connect agents to enterprise tools, databases, and internal systems.

Private Deployment

Run Lyzr locally or on-prem for full data privacy and compliance.

Ready-to-Use Blueprints

Pre-built agent templates for common use cases — deploy instantly.

Multi-Agent & Orchestration

Coordinate agents, automate pipelines, and trigger actions autonomously

Schedulers

Set time-based triggers for autonomous agent execution on a recurring schedule.

Multi-Agent Orchestration

WebSocket Monitoring

Real-time session monitoring — observe agent activity as it happens via WebSocket.

Webhook Triggers

Event-driven execution — agents fire automatically based on external system events.

Audit Log

Full traceability — every agent action is logged for compliance and debugging.

Memory & Context

How Lyzr agents remember, reason, and stay consistent across conversations

Cognis Memory

Stores important user facts and retrieves them at the right moment - ensuring the agent behaves consistently and intelligently across all conversations. No repeated questions. No lost context.

Global Context

Shared context accessible across all agents in a workspace.

Version History

Track changes to agents over time and roll back when needed.

Knowledge Graph

Structured relationships between data entities for smarter retrieval.

Semantic Model

Understands meaning, not just keywords, for precise knowledge retrieval.

Safety & Guardrails

Responsible AI built into the platform

Built-in Hallucination Manager

Automatically detects and prevents hallucinated responses before they reach the user — a core differentiator baked into every Lyrz agent.

Responsible AI Philosophy

Every agent built on Lyrz inherits safety policies by default — hallucination checks, bias filters, and PII protection are on unless explicitly disabled.

Toxicity & Bias Control

Filters harmful, toxic, or biased outputs in real time.

PII Redaction

Automatically detects and redacts personally identifiable information.

Prompt Injection Protection

Blocks malicious prompt injection attacks targeting agent behavior.

Amazon Bedrock Guardrails

Integrates AWS Bedrock's native guardrail policies for enterprise compliance.

LLM as Judge

Uses an independent LLM to validate and score agent responses for accuracy, relevance, and safety — automated quality assurance at scale.

Evaluation & Monitoring

Continuously test, observe, and improve agents in production

AgentEval

Comprehensive agent analytics - track performance, response quality, token usage, and user satisfaction metrics across all deployed agents.

Agent Simulation & Hardening

Systematically test your agent using diverse synthetic user types and edge-case scenarios. Automatically finds weaknesses and improves the agent before production.

Real-time WebSocket Monitoring

Watch agent sessions live — messages, decisions, tool calls, latency — all streamed in real time.

Audit Log

Complete immutable log of every agent action, input, and output — essential for compliance, debugging, and governance.